

mirr0tech for Wildcat — Product Requirements (ETHGlobal Tokyo 2026)

Status	v1, ready to build. Decisions are locked unless marked ⚠️
Owner	Lam (PM)
Team	Lam (PM, pitch, legal content, QA) · Eng A (contracts + chain) · Eng B (backend, integrations, agent) · dashboard shared, see §10
Deadline	Sun 27 Sep 2026, 09:00 JST submission. Internal code freeze Sun 03:00 JST
Repo	LegalMirror/mirr0tech. This PRD extends the existing MVP; it does not replace it
Sponsor tracks	Curvegrid (RWA Tokenization + AI Agent) · Intercepta (Safe Agent-to-Agent Payments) · ENS (Best Use of ENSv2)

Terms: **MLA** = Master Loan Agreement. **AST** = the structured JSON policy we extract from the MLA. **Credential** = Wildcat's on-chain permission for a wallet to deposit into a market. **Borrower** = Wintermute/Selini-style market maker. **Lender** = institution depositing stablecoins.

1. One-sentence pitch

mirr0tech turns a Wildcat Master Loan Agreement into an executable policy. It then screens, credentials and pays institutional lenders on-chain, and every decision traces back to a verbatim clause. Legal teams at borrowers like Wintermute and Selini stop translating contracts into allowlists by hand.

2. Problem

Market makers such as Wintermute and Selini borrow stablecoins from institutional lenders through **Wildcat**, an undercollateralized private-credit protocol. Each market is governed by an off-chain **Master Loan Agreement** (Wildcat publishes a template MLA). The agreement and the chain are only loosely connected:

1. **Manual lender checks.** The borrower's legal/compliance team runs its "Lender Check Process" by hand (KYC/KYB, AML, sanctions, jurisdiction, signed MLA), then manually grants the wallet a deposit credential on-chain.
2. **Legal-to-code translation risk.** The link between an MLA clause and the on-chain allowlist lives in someone's head and spreadsheet. Nothing proves which clause justified which on-chain action.
3. **Point-in-time screening.** A lender screened at onboarding can become sanctioned later. The MLA says what must happen then; enforcing it is manual.
4. **Payout risk.** Interest and principal repayments go out to lender wallets without a fresh screen at the moment of payment.
5. **Audit pain.** Reconstructing "why was this wallet allowed / paid / blocked" for regulators or counterparties takes days.

Who feels it: the borrower's legal & compliance team (primary), the borrower's treasury/ops (secondary), and institutional lenders waiting to be onboarded (tertiary).

⚠️ **Lam to confirm before the pitch:** that Wintermute and Selini have in fact run Wildcat markets, and a rough "time to onboard one lender today" figure (hours or days). That figure is the pitch's headline metric. Present them as **target customers**, never as partners, unless they have agreed to it.

3. Goals, non-goals, success criteria

Goals (hackathon)

- G1. Ingest a real-shaped MLA (Wildcat template + illustrative term sheet) and produce a source-quoted policy AST.
- G2. Onboard a lender end-to-end on **Sepolia**: live Intercepta screen → policy decision → ENSv2 subname → Wildcat-style credential, with **hybrid** human approval.
- G3. A **treasury agent** that pays lenders and blocks or holds a payout when Intercepta flags the recipient, showing its reasoning.
- G4. A compliance **dashboard** showing policy ↔ clause, approval queue, lenders, payouts and the audit stream.
- G5. Meet every hard requirement of the three sponsor tracks (§9).

Non-goals (explicitly out of scope)

- Real Wildcat V2 contracts or mainnet. We ship a **minimal Wildcat-style mock** (§6.2).
- Real KYC/KYB providers, real fiat or real funds. Mock USDC only.
- Investor/lender self-service auth. The API stays an **operator API**.
- PDF/OCR ingestion, long-document chunking, legal correctness guarantees.
- Interest accrual math, penalty APR, delinquency and default state machine (named in the pitch as roadmap).

Demo success criteria (must all be true at freeze)

- ☐ Golden path in §4 runs on Sepolia from the dashboard in < 4 minutes, with no terminal.
- ☐ At least one **denied/held** path for onboarding and for payout is shown live, with the clause quote.
- ☐ Every sponsor checklist item in §9 is ticked.
- ☐ Public GitHub repo, README complete, demo video uploaded, live demo URL works.

4. Demo script (the golden path)

This is the spec. If a feature doesn't show up in this script, it's P1 or lower.


#	Screen	What happens	Sponsor shown
1	Policy	Officer opens the MLA (Wildcat template + "Borrower: Demo MM Ltd" term sheet). Clicks Compile . Rules appear, each with its clause number and highlighted quote. Policy hash shown. Optional: Re-extract with LLM button shows live extraction	—
2	Policy	Click Deploy market . Market + hook deployed/registered through MultiBaas with the policy hash stored on-chain	Curvegrid
3	Lenders	Add <i>Lender A</i> (clean mainnet address). Gateway calls Intercepta live → clean. Policy: all facts true → auto-approved	Intercepta
4	Lenders	Gateway issues <code>lender-a.credit.<borrower>.eth</code> on ENSv2 Sepolia with text records <code>mirr0tech.status=approved</code> , <code>mirr0tech.policyHash</code> , <code>mirr0tech.credentialExpiry</code> . Then grants the Wildcat-style credential on-chain	ENS, Curvegrid
5	Lenders	Lender A deposits 10,000 mUSDC into the market → succeeds. Event appears via MultiBaas webhook	Curvegrid
6	Queue	Add <i>Lender B</i> with a missing fact (e.g. <code>mlaSigned</code> unknown). Policy returns needs review . Officer sees the trace, clicks Approve after ticking the attestation → credential issued	—
7	Lenders	Add <i>Lender C</i> : a real sanctioned mainnet address . Intercepta flags it → policy <code>forbid</code> → denied, no override (hard rule). Deposit attempt from C's wallet reverts on-chain	Intercepta

#	Screen	What happens	Sponsor shown
8	Agent	Treasury agent runs the interest payout cycle: reads market state via MultiBaas, plans payouts to A and B, re-screens each recipient with Intercepta before signing .	Intercepta, Curvegrid AI
9	Agent	Pays A ✓. B's address has since been flagged (demo: switched to a flagged address, see §12) → payout held , credential revoked, ENS status → suspended , explanation cites the MLA's sanctions clause	Intercepta, ENS
10	Audit	Timeline of every decision with rule trace + clause quote + tx hash. Close with: "every on-chain action is traceable to a sentence in the contract."	—

5. Users & key user stories

As a...	I want to...	So that...	Priority
Compliance officer	upload an MLA and see each executable rule next to the clause it came from	I can check the machine's reading before anything goes live	P0
Compliance officer	have clean lenders auto-approved and only exceptions queued for me	my time goes to edge cases	P0
Compliance officer	never be able to approve a sanctioned wallet	a mis-click can't become a legal breach	P0
Treasury ops	have payouts screened at the moment of payment	we never pay a newly sanctioned wallet	P0
Lender	get a human-readable identity (<code>lender-a.credit.<borrower>.eth</code>) showing my status	counterparties can verify me without seeing my KYC data	P0
Auditor / counterparty	see why any wallet was allowed, paid or blocked	we can answer regulators in minutes	P0
Compliance officer	revoke a lender and have credential + ENS status update together	state never drifts across systems	P1

6. Solution & architecture



6.1 What we keep from the current repo (do not rewrite)

- Document normalization, SHA-256 provenance and verbatim-quote validation (`src/policy/document.js`, `schema.js`).
- Strict JSON-schema AST, fail-closed evaluator with rule traces (`evaluate.js`).
- LLM extraction via OpenAI Structured Outputs (`extract.js`), no silent fallback.

- Idempotency keys, operation intents before signing, restart recovery, audit log (`service.js`, `store.js`).
- Security stance: **never execute LLM-written code**; the LLM only emits data validated against the schema.

6.2 What changes

Area	Today	Change
Domain	RWA fund token (mint/burn investors)	Wildcat credit market (lenders, credentials, payouts)
AST action enum	<code>mint</code> , <code>burn</code> , <code>transfer</code>	<code>credential</code> , <code>payout</code> , <code>deposit</code> (keep old values out)
FACTS	RWA facts	See §7.1
Contract	<code>MirrorToken.sol</code> (custodial ERC-20)	<code>MockWildcatMarket.sol</code> + <code>MirrorAccessHook.sol</code> + <code>MockUSDC.sol</code> (keep <code>MirrorToken</code> in repo, unused)
Chain lock	chain ID 31337 only	Allow Sepolia (11155111) + 31337 for tests
Chain calls	ethers directly	Via MultiBaas REST for deploy, calls, events. Keep ethers for tests
Demo doc	Securitize/BlackRock services agreement	Wildcat template MLA + illustrative term sheet
UI	none	Dashboard (§7.6)

7. Functional requirements

Priority: **P0** = required for the demo · **P1** = do if P0 is done by Sat 18:00 · **P2** = stretch / pitch-only.

7.1 Policy & legal content (owner: Lam + Eng B)

- **P0** Prepare `test/legal/wildcat-mla.md`: the Wildcat template MLA plus an **illustrative** Exhibit A term sheet (borrower "Demo MM Ltd", asset mUSDC, base APR, lender check process description). No real counterparty names in the document itself. ⚠ Check the template's license/terms before committing it to a public repo; if unclear, link to it and commit only excerpts.
- **P0** New fact set (booleans; unknown = fail closed):

Fact	Source	Used by
<code>sanctionsClear</code>	Intercepta (live)	<code>credential</code> , <code>deposit</code> , <code>payout</code>
<code>amlApproved</code>	Intercepta risk grade below threshold (live)	<code>credential</code> , <code>payout</code>
<code>kycApproved</code>	Operator attestation (mock)	<code>credential</code>
<code>mlaSigned</code>	Operator attestation (mock)	<code>credential</code>
<code>lenderCheckPassed</code>	Operator attestation per the MLA's Lender Check Process	<code>credential</code>
<code>credentialActive</code>	On-chain hook read	<code>deposit</code> , <code>payout</code>
<code>reservesSufficient</code>	On-chain market read by the agent	<code>payout</code>

- **P0** Hand-authored fixture AST for the MLA (same pattern as `src/policy/fixture.js`): ~6–8 rules, each quoting the MLA verbatim. Must include a `forbid` rule on `not sanctionsClear` for **both** `credential` and `payout`.
- **P0** Everything the MLA leaves undecided goes in `unresolved` (e.g., jurisdiction eligibility isn't a boolean the MLA defines; Process Agent appointment; oracle-error risk allocation).
- **P1** Live extraction on the MLA via the existing OpenAI path, shown as a side-by-side "LLM candidate vs reviewed fixture" diff. The demo must still run on the fixture if the API is slow.
- **P1** Clause highlighting: `GET /v1/policy` returns character offsets for each quote so the UI can highlight it.

7.2 Decision engine & hybrid automation (owner: Eng B)

Evaluation result maps to three outcomes:

Outcome	Condition	What happens
approve	every <code>require</code> satisfied, a <code>permit</code> matches, no <code>forbid</code> matches, and all facts known	auto-execute: ENS subname + credential grant
review	no <code>forbid</code> matches, but at least one fact is unknown or an attestation is missing	into the approval queue; the officer sets the missing attestations, then re-evaluate
deny	any <code>forbid</code> matches (e.g., sanctions)	hard stop. No officer override . Audit + dashboard alert

- **P0** Extend `evaluate.js` to return `approve` | `review` | `deny` (today it's permit/deny with trace). Unknown facts must never produce `approve`.
- **P0** Re-evaluation is idempotent and fully traced in the audit log.

7.3 Contracts (owner: Eng A)

Model on Wildcat V2's **access-control hooks** / **role-provider** pattern, kept minimal. ⚠ We're copying the concept, not the interface; label it "Wildcat-style" in README and pitch.

- **P0** `MockUSDC.sol`: 6 decimals, open `mint` for the demo.
- **P0** `MirrorAccessHook.sol`
- `bytes32 immutable policyHash, PROVIDER_ROLE (gateway wallet), DEFAULT_ADMIN_ROLE`.
- `grantCredential(address lender, uint64 expiry, bytes32 operationId) / revokeCredential(address lender, bytes32 operationId);` operation IDs unique (reuse the processed pattern).
- `isCredentialed(address)` view (active and not expired).
- Events: `CredentialGranted`, `CredentialRevoked` (include `policyHash`).
- **P0** `MockWildcatMarket.sol`
- `deposit(uint256)` reverts unless `hook.isCredentialed(msg.sender)`; tracks lender balances.
- `payout(address lender, uint256 amount, bytes32 operationId)` callable only by `TREASURY_ROLE` (agent wallet) and only if `isCredentialed(lender)`.
- `totalAssets`, `balanceOf(lender)` views for the agent.
- **P0** Foundry/Node tests: credentialed deposit ok, uncredentialed reverts, revoked blocks payout, expiry, duplicate op ID, roles.
- **P0** Deploy script for Sepolia; addresses to `generated/deployment.json`; **verified on Etherscan**.
- **P1** Withdrawal request flow (open-term) for realism.

7.4 Integrations (owner: Eng B, Eng A for MultiBaas deploy)

Intercepta (P0) - `src/integrations/intercepta.js`: one function `screen(address) → { sanctionsClear, amlApproved, grade, reasons[], raw, checkedAt }`. - Called **live** on: lender onboarding, and **before every payout is signed** (hard requirement). - Screen **real mainnet addresses**, even though payments run on Sepolia (track requirement). Keep a demo list: 2 clean well-known addresses and 1 known-sanctioned address, each mapped to a Sepolia demo wallet. - Timeout/unavailable → facts become **unknown** → `review` (onboarding) or **held** (payout). Never approve on error. - Store the raw response in the audit record. Get the free key at intercepta.io/ethglobal.

ENSv2 on Sepolia (P0) - Borrower parent name, e.g. `credit.<team-name>.eth` (use a name we own; do not register or imply `wintermute.eth`). - On approve: create subname `<lender-slug>.credit.<...>.eth` → lender address; set text records `mirr0tech.status` (`approved|suspended`), `mirr0tech.policyHash`, `mirr0tech.credentialExpiry`, `mirr0tech.auditUrl`. - Use ENSv2 per-record permissions (Permissioned Resolver) so the gateway wallet can edit **only** the `mirr0tech.*` records. That delegation is what makes ENS "central, not cosmetic"; show it in the pitch. - Subnames should be non-transferable by the lender if ENSv2 supports that; otherwise document it as a known limitation. - ⚠ ENSv2 is a beta and its Sepolia state has

been reset before. Eng B does a **spike first (Fri night, ≤2h)**: register parent, create a subname, set a text record from a script. If it fails, fall back to ENS v1 on Sepolia and tell Lam immediately (this affects eligibility for the ENS track). - Dashboard shows the resolved name, never raw addresses where a name exists.

Curvegrid MultiBaas (P0) - Create a MultiBaas deployment on Sepolia; upload the 3 contract ABIs; link deployed addresses. - Gateway calls grant/revoke/payout through MultiBaas REST (signing via MultiBaas or local wallet + MultiBaas tx submission, whichever is faster to set up; document which). - **Event webhooks** (CredentialGranted, Deposit, Payout) → POST /v1/webhooks/multibaas → audit stream. - README section "How we use MultiBaas" + feedback (required).

7.5 Treasury agent (owner: Eng B, P0 core / P1 LLM)

Goal: a **policy-aware transaction agent** (Curvegrid AI Agent track) whose payments are screened (Intercepta track).

- **P0** src/agent/treasury.js, triggered by POST /v1/agent/run (dashboard button) and optionally on an interval.
- Loop per cycle: 1. Read market state via MultiBaas (lender balances, reserves). 2. Plan payouts (demo: fixed interest amount per lender, pro-rata). 3. For each recipient: **Intercepta screen → facts → policy payout evaluation**. 4. approve → sign & send payout → record tx. deny/unknown → **hold**, revoke credential, set ENS status suspended, raise alert. 5. Emit a decision record: { lender, ens, amount, decision, failedRules, clauseQuotes, interceptaReasons, txHash? }.
- **P0** The decision is deterministic (policy engine). **P1** An LLM writes the human-readable explanation from the decision record (read-only; it cannot change the decision).
- **P2** x402: expose GET /v1/agent/report behind an x402 paywall, or have the lender-side agent pay a deposit via x402. The track says x402 is *preferred*, not required.
- Agent wallet holds TREASURY_ROLE only. It cannot grant credentials.

7.6 Dashboard (owner: shared; Lam owns copy and flow)

Stack: whatever is fastest for the team (suggest Next.js + Tailwind, or plain Vite). Calls the gateway via a server-side proxy so the API key never reaches the browser.

- **P0 Policy**: document viewer with highlighted quotes; rule list (action, effect, condition, clause, rationale); unresolved terms; policy hash; Compile / Deploy buttons.
- **P0 Lenders**: add lender (name, address, demo mainnet address for screening); status chips (approved / review / denied / suspended); ENS name; credential expiry; Intercepta grade + reasons.
- **P0 Queue**: pending reviews with the trace ("missing: mlaSigned — clause 3(j) ..."), attestation checkboxes, Approve / Reject.
- **P0 Agent**: "Run payout cycle" button; live list of planned → paid / held payouts with explanations and Etherscan links.
- **P0 Audit**: reverse-chronological stream, filter by lender, each row expandable to trace + raw Intercepta response + tx.
- **P1** Wildcat-style market summary card (total deposits, lenders, reserves). This also helps with the Curvegrid "Digital Asset Dashboard" framing.

7.7 API changes (owner: Eng B)

Keep the conventions: decimal-string amounts, Idempotency-Key on state changes, operator bearer auth, typed errors.

Method	Endpoint	Purpose	Pri
GET	/v1/policy	Compiled policy + quotes (+ offsets P1)	P0
POST	/v1/policy/compile	Compile fixture or uploaded candidate	P1
POST	/v1/lenders	{ name, address, screeningAddress? } → runs screen + evaluation	P0

Method	Endpoint	Purpose	Pri
GET	/v1/lenders[:id]	Lender, facts, decision, ENS, credential	P0
PATCH	/v1/lenders/:id/attestations	Set kycApproved, mlaSigned, lenderCheckPassed → re-evaluate	P0
POST	/v1/lenders/:id/approve /reject	Resolve a review item (rejected if any forbid matches)	P0
POST	/v1/lenders/:id/revoke	Revoke credential + ENS suspended	P1
POST	/v1/lenders/:id/rescreen	Fresh Intercepta call	P1
POST	/v1/agent/run	Run a payout cycle	P0
GET	/v1/payouts[:id]	Payout decisions + tx	P0
POST	/v1/webhooks/multibaas	Event ingestion (verify MultiBaas signature)	P0
GET	/v1/audit	Audit stream	P0

Old investors/mints/redemptions/deposits endpoints: remove or leave unused. Don't spend time migrating tests for them; delete those tests if they block CI.

8. Non-functional & security requirements

- **Fail closed** everywhere: unknown fact, API timeout, RPC uncertainty → no approval, no payout.
- **No LLM code execution**; LLM output is schema-validated data only (existing guarantee; keep the tests).
- **Secrets** (API_KEY, Intercepta, OpenAI, MultiBaas, private keys) only in .env; nothing in git; dashboard never sees the API key.
- **Separate wallets**: deployer/admin, gateway (PROVIDER_ROLE), agent (TREASURY_ROLE). Testnet keys only.
- **Privacy**: no KYC data on-chain or in ENS records, only status, policy hash, expiry.
- **Audit integrity**: the Gemini draft calls the audit log "unalterable". **It isn't**: today it's an editable local JSON file. Say "tamper-evident" only if we do P2: hash-chain the audit entries and anchor the head hash on-chain or in an ENS text record.
- **Disclaimers** in README and UI footer: prototype, testnet, not legal advice, no affiliation with Wildcat, Wintermute or Selini.

9. Sponsor requirement checklists

⚠ Taken from the ETHGlobal prize page summary on 25 Sep. **Lam re-reads the live prize pages Saturday morning** and updates this section. Also confirm how many sponsor prizes one project may select.

Curvegrid: Best RWA Tokenization Project + Best AI Agent Project (\$1k each) - ☐ GitHub repo with contracts, tests, documentation, solid README - ☐ README: one-sentence summary · how MultiBaas is used · team intro with social handles · setup & test instructions · MultiBaas feedback - ☐ Agent understands blockchain activity (reads market state) and takes on-chain action (payouts, revocations)

Intercepta: Safe Agent-to-Agent Payments with x402 (\$1,250 / \$750) - ☐ Working agent payment flow (testnet OK; x402 preferred → \$7.5 P2) - ☐ ≥1 **live** Intercepta API call runs **before a payment is signed**, and its result decides the outcome - ☐ Screens **real mainnet addresses** - ☐ Demo shows one payment approved and one blocked/held, with visible reasoning - ☐ README points to the files that call the API + 3-5 lines of API feedback

ENS: Best Use of ENSv2 (\$3k / \$2k / \$1k) - ☐ Built on ENSv2 beta (Sepolia) - ☐ ENSv2 is central: subname hierarchy per borrower + per-record delegated permissions + status records read back by the dashboard/agent - ☐ Functional, not hard-coded values - ☐ Live demo link + open-source GitHub - ☐ Bonus: agent integration (agent updates ENS status on hold/suspend)

ETHGlobal general - ☐ Meaningful commit history across the weekend (no single final-day dump) - ☐ Demo video ($\leq$ the event's limit; confirm), project description, screenshots - ☐ Live URL for dashboard (e.g. Vercel + gateway on Render/Fly/Railway)

10. Workstreams, owners, timeline

Times are **JST**. Now $\approx$ Fri 18:00. $\sim$ 39h to deadline, $\sim$ 33h to freeze.

Workstream	Owner	P0 scope
WS1 Legal & policy	Lam (content) + Eng B (code)	MLA doc, facts, fixture AST, 3-outcome evaluator
WS2 Contracts & chain	Eng A	Hook, market, mUSDC, tests, Sepolia deploy + verify, MultiBaas setup
WS3 Integrations	Eng B	Intercepta, ENSv2, MultiBaas calls + webhooks
WS4 Treasury agent	Eng B (Eng A helps Sat pm)	Payout loop, hold logic, explanations
WS5 Dashboard	Eng A (Sat pm) + Lam	5 screens from \$7.6
WS6 Pitch & submission	Lam	Script, video, README, sponsor forms, feedback text

⚠ Lam: assign names to Eng A / Eng B. If one engineer is stronger on frontend, swap WS5 to them and move WS4 to Eng A.

When (JST)	Milestone	Exit check
Fri 22:00	Spikes done: ENSv2 subname + text record from script; Intercepta key works on 1 clean + 1 sanctioned address; MultiBaas deployment created	Each spike has a script committed in <code>scripts/spikes/</code>
Sat 02:00	Contracts + tests pass locally; MLA doc + fixture AST committed; fact set + 3-outcome evaluator merged	<code>npm test</code> green
Sat 12:00	Contracts on Sepolia (verified) and linked in MultiBaas; onboarding path (steps 3-4, 6-7) works via API/curl	Lender A approved end-to-end on Sepolia
Sat 18:00	Agent payout approve + hold working (steps 8-9); dashboard skeleton reads live API	Full golden path via API
Sat 24:00	Dashboard P0 complete; deployed to public URL; P1s only if green	Lam runs golden path from UI 2x
Sun 03:00	Code freeze. Only bug fixes after this	Tag <code>v0.1-freeze</code>
Sun 03:00-07:00	Record demo video, finish README, sponsor feedback sections, screenshots	Video uploaded
Sun 08:00	Submit (1h buffer)	ETHGlobal submission confirmed

Rules of engagement: small PRs to `main`, commit often (history matters), any P0 slipping $> 2h \rightarrow$ tell Lam, cut scope using §11.

11. Scope cut order (if behind)

Cut from the top first: 1. x402 (P2) · tamper-evident audit (P2) 2. LLM explanation for the agent (P1). Use templated text built from the trace 3. Live LLM extraction in the demo (P1). Show the fixture only; mention extraction exists in the repo 4. Withdrawals, revoke/rescreen endpoints (P1) 5. MultiBaas webhooks. Poll events via MultiBaas REST instead 6. **Never cut:** live Intercepta call before payout, one approved + one blocked, ENSv2 subname + records, Sepolia deploy, source-quoted rules

12. Risks & mitigations

Risk	Likelihood	Mitigation
ENSv2 Sepolia unstable/reset or missing features	Med	Fri-night spike; fallback ENS v1 Sepolia; record a backup video of the working flow early
Intercepta sanctioned-address demo: payout target changes status mid-demo	Certain (by design)	Payout recipient for Lender B maps to a different screening address at payout time via <code>screeningAddress</code> override; disclose this in the demo ("simulating a status change")
Intercepta rate limits / downtime during judging	Med	Fail-closed path is itself a valid demo; cache last result for display only, never for decisions
MultiBaas setup takes longer than expected	Med	Eng A starts it Fri night; ethers fallback exists in repo; README still documents the attempt
LLM extraction produces bad rules live	Med	Fixture is the default path; extraction only shown as a diff
Legal/IP: MLA template license, use of firm names	Low-Med	Check license; illustrative term sheet; "target customers" wording; disclaimers
Scope too big for 3 people / 33h	High	Cut order §11; golden path first, polish last

13. Decisions log (vs Gemini draft)

Topic	Gemini draft	Decision	Why
On-chain target	Wildcat <code>MarketController</code> whitelist hook	Minimal Wildcat-style market + access hook	Wildcat docs list V2 as mainnet-only; the Sepolia deployment is an incomplete 2024 build
Intercepta usage	Onboarding screen	Onboarding and a treasury agent screening each payout before signing	Track requires an agent payment flow with approve + block
AI agent	Autonomous treasury agent (open question)	Yes, narrow: payout agent; the decision is deterministic, the LLM explains	Serves Intercepta + Curvegrid AI + ENS bonus with one feature
Automation	Open question	Hybrid : auto-approve clean, queue unknowns, hard-deny forbids	Legal teams keep control of exceptions; shows both paths
Legal ingestion	Open question	Reviewed fixture by default; live OpenAI extraction as a side-by-side	Demo reliability; extraction already implemented
Third sponsor	ENSv2	ENSv2 (kept)	Team choice; World ID for Agents noted as future work
Audit log	"Unalterable"	"Local audit log"; tamper-evident only if P2 done	Accuracy in front of judges
Wildcat link in draft	Pointed to an ASX mining company	Use docs.wildcat.finance	Wrong link

14. Open questions

1. ⚠️ Owner names for Eng A / Eng B, and each person's strongest stack. (Lam, Fri 19:00)
2. ⚠️ Do Wintermute / Selini run Wildcat markets today, and what's the current per-lender onboarding time? (Lam, before the pitch)
3. Which parent ENS name do we register on Sepolia? (Eng B during the spike)
4. MultiBaas signing: MultiBaas-managed signer (HSM wallet) or local key + submit? (Eng A, Fri night)
5. Hosting for the gateway (Render / Fly / Railway) and dashboard (Vercel). (Eng A, Sat)

6. Demo video length limit and number of selectable sponsor prizes. (Lam, Sat am)

15. References

- Wildcat docs: [Hooks](#) · [Market access via policies/hooks](#) · [Contract deployments](#) · [Template MLA](#) · [V2 audit repo](#)
- ENS: [ENSv2 overview](#) · [ENSv2 beta announcement](#)
- Curvegrid: [MultiBaas docs](#) · [Sample app](#)
- Intercepta: [Compliance Engine](#) · free hackathon key at intercepta.io/ethglobal
- ETHGlobal Tokyo 2026 [prizes](#)
- Repo README (current MVP architecture, API conventions, recovery model)